

Multimodal Property RAG System – Technical Analysis Report

Executive Summary

This technical report analyzes the **Multimodal Retrieval-Augmented Generation (RAG)** system designed for property data analytics, focusing on the **dual analytics engine integration** using **Google Gemini** and **OpenAI ChatGPT**, and diagnosing the issue where ChatGPT analytics results are not displayed in the Flask frontend.

Key Findings

- The system successfully integrates **Gemini** and **ChatGPT** for parallel pandas code generation and execution.
 - **ChatGPT analytics function correctly** in backend testing but fail to appear in the Flask frontend.
 - **Root cause identified:** Missing OpenAI API key in the production environment.
 - The architecture demonstrates **robust fault tolerance**, **safe code execution**, and **strong fallback mechanisms**.
-

1. System Architecture Overview

1.1 Core Components

Vector Database Layer

- **Technology:** Pinecone (Serverless)
- **Embeddings:** HuggingFace SentenceTransformer ([all-MiniLM-L6-v2](#))
- **Dimension:** 384

- **Records:** 147,666 property vectors
- **Index Name:** `property-data-rag`

Analytics Layer (Dual Engine)

- **Gemini Analytics Agent:** Google Gemini 2.5 Flash
- **ChatGPT Analytics Agent:** OpenAI GPT-4o-mini
- **Function:** Natural language → Pandas code translation
- **Execution:** Safe code execution with strict heuristics and keyword filtering

Response Generation Layer

- **Primary Composer:** GPT-4o-mini
- **Fallback Composer:** Gemini 2.5 Flash
- **Input Context:** Combined output from vector search and both analytics agents

Frontend Layer

- **Framework:** Flask
- **UI Stack:** HTML, CSS, JavaScript (ES6)
- **Theme:** Dark, responsive design with toggle-based output panels
- **Runtime:** AJAX-driven request–response cycle

2. Data Pipeline

- Property CSV → EDA Processing → Cleaned Data → Embeddings → Pinecone Index
- ↓
- Analytics DataFrame → Gemini & ChatGPT Agents → Pandas Code Execution
- ↓

- User Query → Vector Search + Dual Analytics → Context Assembly → AI Response
-

3. Technical Implementation Analysis

3.1 Multimodal Analytics Implementation

Gemini Analytics Agent

- `class DataAnalyticsAgent:`
- `def analyze(self, nl_query: str, top_k: int = 7) -> Dict[str, Any]:`
- `# Uses Gemini 2.5 Flash for pandas code generation`
- `# Implements safety filters and fallback heuristics`
- `return {"context": str, "rows": list, "generated_code": str}`

Strengths

- Effective fallback heuristics
- Secure execution filters
- Reliable code extraction and structured output

ChatGPT Analytics Agent

- `class ChatGPTAnalyticsAgent:`
- `def analyze(self, nl_query: str, top_k: int = 7) -> Dict[str, Any]:`
- `# Uses OpenAI GPT-4o-mini for pandas code generation`
- `# Implements safety mechanisms and markdown code extraction`

Highlights

- Dual API strategy (OpenAI SDK + LangChain fallback)
- Enhanced logging for debugging

- Identical prompt structure to Gemini for consistency

3.2 Dual Execution Architecture

- `def analyze_both(self, query: str, top_k: int = 7) -> Dict[str, Any]:`
- `"""Run both Gemini and ChatGPT analytics and return both results."""`
- `gem = self.analytics_agent.analyze(query, top_k=top_k)`
- `chg = self.chatgpt_analytics_agent.analyze(query, top_k=top_k)`
- `return {"gemini": gem, "chatgpt": chg}`

- **Design Pattern:** Currently sequential; parallel execution planned for future optimization.

3.3 Frontend Integration

Flask API Endpoint

- `@app.route('/api/ask', methods=['POST'])`
- `def api_ask():`
- `both_analytics = RAG_AGENT.analyze_both(q, top_k=7)`
- `return jsonify({`
- `"answer": str(answer),`
- `"analytics": both_analytics,`
- `"matches": matches_serializable`
- `})`

JavaScript Frontend Rendering

- `const gemCode = j.analytics?.gemini?.generated_code || "";`
 - `const chgCode = j.analytics?.chatgpt?.generated_code || "";`
 - `document.getElementById('chatgpt-code-content').textContent =`
 - `chgCode || 'No pandas code generated';`
-

4. Issue Analysis: ChatGPT Analytics Not Displaying

4.1 Symptom

- ChatGPT panel shows “No pandas code generated.”
- Backend tests confirm ChatGPT analytics function correctly.
- Issue isolated to Flask production runtime.

4.2 Root Cause: Missing OpenAI API Key

- OPENAI_API_KEY found: False

Without this key, the OpenAI client fails silently, resulting in empty analytics responses.

4.3 Secondary Causes

1. Improper `.env` loading due to different working directories.
2. Missing error reporting for failed API initialization.
3. Caching of invalid client state (`None`) on startup.

4.4 Code Flow Comparison

Working Path (Backend Test)

- `test_rag_system.py` → `pinecone_rag_setup.py` → `ChatGPTAnalyticsAgent.analyze()`
- → `get_openai_client()` → OpenAI API → Returns valid pandas code

Failing Path (Flask Frontend)

- `app.py` → `RAG_AGENT.analyze_both()` → `ChatGPTAnalyticsAgent.analyze()`

- → `get_openai_client()` → None (API key missing) → Empty output
-

5. Technical Recommendations

5.1 Immediate Fixes

Environment Configuration

- `# .env file`
- `OPENAI_API_KEY=sk-proj-your-key-here`
- `PINECONE_API_KEY=your-pinecone-key`
- `GEMINI_API_KEY=your-gemini-key`

Improved Client Initialization

- `def get_openai_client():`
- `client = setup_openai()`
- `if not client:`
- `print("[ERROR] OpenAI client not initialized - missing API key")`
- `return client`

Frontend Error Display

- `if (!chgCode) {`
 - `document.getElementById('chatgpt-code-content').textContent =`
 - `'ChatGPT analytics unavailable - check API configuration';`
 - `}`
-

5.2 Architecture Enhancements

Parallel Execution

- `import asyncio`
- `from concurrent.futures import ThreadPoolExecutor`

-
- `async def analyze_both_parallel(self, query: str, top_k: int = 7):`
- `with ThreadPoolExecutor(max_workers=2) as executor:`
- `gem = executor.submit(self.analytics_agent.analyze, query, top_k)`
- `chg = executor.submit(self.chatgpt_analytics_agent.analyze, query, top_k)`
- `return {"gemini": gem.result(), "chatgpt": chg.result()}`

Health Check Endpoint

- `@app.route('/api/health')`
- `def health_check():`
- `return jsonify({`
- `"pinecone": bool(RAG_AGENT.index),`
- `"gemini": bool(RAG_AGENT.model),`
- `"openai": bool(get_openai_client())`
- `})`

5.3 Performance Optimizations

Caching Analytics

- `from functools import lru_cache`
-
- `@lru_cache(maxsize=1000)`
- `def cached_analytics(query_hash, agent):`
- `# Store recent query results to minimize API calls`
- `pass`

Batch Query Processing

- `def batch_analyze(self, queries: List[str], top_k: int = 7):`
 - `return [self.analyze_both(q, top_k) for q in queries]`
-

6. Security Considerations

6.1 API Key Management

- Use secret management (AWS Secrets Manager, Azure Key Vault).
- Avoid storing credentials in `.env` in production.
- Implement API key rotation policies.

6.2 Safe Code Execution

- `forbidden = ['os.', 'sys.', 'eval(', 'exec(', 'import ', '___']`
- for token in forbidden:
- `code = code.replace(token, '# removed ')`

6.3 Input Validation

- `def validate_query(q):`
 - `return len(q) < 500 and "drop" not in q.lower()`
-

7. Testing Strategy

Unit Tests

- `def test_chatgpt_analytics():`
- `agent = ChatGPTAnalyticsAgent(test_df)`
- `res = agent.analyze("2 bedroom under £2000")`
- `assert res["generated_code"]`

Integration Tests

- `def test_multimodal_rag():`
- `r = PropertyRAGAgent().analyze_both("test query")`
- `assert "gemini" in r and "chatgpt" in r`

Load Tests

Simulate 10 concurrent queries using threading to measure latency.

8. Deployment & Monitoring

Deployment Checklist

1. `.env` configured with all keys
2. Verify individual services (Pinecone, Gemini, OpenAI)
3. Run integration tests
4. Enable logging and monitoring

Logging Example

- `logging.info(f"[ANALYTICS] ChatGPT | {duration:.2f}s | {query[:60]}")`

Error Recovery

- if not openai_available:
 - `return {"chatgpt": {"context": "Analytics unavailable", "rows": [], "generated_code": ""}}`
-

9. Conclusion

The **Multimodal Property RAG System** showcases a robust architecture combining **semantic retrieval** and **AI-driven data analytics**. The dual-engine design enhances resilience and analytical diversity, offering improved reliability across varied data contexts.

Key Achievements

- Dual LLM integration (Gemini + ChatGPT)
- Intelligent fallback and error handling
- Safe code execution framework
- Real-time analytics UI

Required Actions

1. Add missing OpenAI API key.
2. Implement real-time health checks.
3. Introduce parallel execution and caching.

Once configured, the system is **production-ready** and capable of scaling for enterprise-grade property analytics.